

# CO2 Prediction Analysis

Shubhan Tamhane

2025-01-18

Including Plots

Importing necessary libraries

```
library(dplyr)
```

```
##  
## Attaching package: 'dplyr'  
  
## The following objects are masked from 'package:stats':  
##  
##   filter, lag  
  
## The following objects are masked from 'package:base':  
##  
##   intersect, setdiff, setequal, union
```

```
library(car)
```

```
## Loading required package: carData  
  
##  
## Attaching package: 'car'  
  
## The following object is masked from 'package:dplyr':  
##  
##   recode
```

```
library(randomForest)
```

```
## randomForest 4.7-1.2  
  
## Type rfNews() to see new features/changes/bug fixes.  
  
##  
## Attaching package: 'randomForest'  
  
## The following object is masked from 'package:dplyr':  
##  
##   combine
```

```
library(ggplot2)
```

```
##  
## Attaching package: 'ggplot2'  
  
## The following object is masked from 'package:randomForest':  
##  
##     margin
```

Importing data

```
d1 <- read.csv("/Users/shubhantamhane/Downloads/CO2 Emissions.csv")  
colnames(d1)
```

```
## [1] "Make" "Model"  
## [3] "Vehicle.Class" "Engine.Size.L."  
## [5] "Cylinders" "Transmission"  
## [7] "Fuel.Type" "Fuel.Consumption.City..L.100.km."  
## [9] "Fuel.Consumption.Hwy..L.100.km." "Fuel.Consumption.Comb..L.100.km."  
## [11] "Fuel.Consumption.Comb..mpg." "CO2.Emissions.g.km."
```

Checking for null values

```
sum(is.na.data.frame(d1))
```

```
## [1] 0
```

Exploring data

```
head(d1)
```

```
##      Make      Model Vehicle.Class Engine.Size.L. Cylinders Transmission  
## 1 ACURA      ILX      COMPACT      2.0          4          AS5  
## 2 ACURA      ILX      COMPACT      2.4          4          M6  
## 3 ACURA ILX HYBRID      COMPACT      1.5          4          AV7  
## 4 ACURA      MDX 4WD    SUV - SMALL      3.5          6          AS6  
## 5 ACURA      RDX AWD    SUV - SMALL      3.5          6          AS6  
## 6 ACURA      RLX      MID-SIZE      3.5          6          AS6  
##      Fuel.Type Fuel.Consumption.City..L.100.km. Fuel.Consumption.Hwy..L.100.km.  
## 1          Z          9.9          6.7  
## 2          Z         11.2          7.7  
## 3          Z          6.0          5.8  
## 4          Z         12.7          9.1  
## 5          Z         12.1          8.7  
## 6          Z         11.9          7.7  
##      Fuel.Consumption.Comb..L.100.km. Fuel.Consumption.Comb..mpg.  
## 1          8.5          33  
## 2          9.6          29  
## 3          5.9          48  
## 4         11.1          25
```

```
## 5          10.6          27
## 6          10.0          28
## CO2.Emissions.g.km.
## 1          196
## 2          221
## 3          136
## 4          255
## 5          244
## 6          230
```

```
str(d1)
```

```
## 'data.frame': 7385 obs. of 12 variables:
## $ Make : chr "ACURA" "ACURA" "ACURA" "ACURA" ...
## $ Model : chr "ILX" "ILX" "ILX HYBRID" "MDX 4WD" ...
## $ Vehicle.Class : chr "COMPACT" "COMPACT" "COMPACT" "SUV - SMALL" ...
## $ Engine.Size.L. : num 2 2.4 1.5 3.5 3.5 3.5 3.5 3.7 3.7 2.4 ...
## $ Cylinders : int 4 4 4 6 6 6 6 6 6 4 ...
## $ Transmission : chr "AS5" "M6" "AV7" "AS6" ...
## $ Fuel.Type : chr "Z" "Z" "Z" "Z" ...
## $ Fuel.Consumption.City..L.100.km.: num 9.9 11.2 6 12.7 12.1 11.9 11.8 12.8 13.4 10.6 ...
## $ Fuel.Consumption.Hwy..L.100.km. : num 6.7 7.7 5.8 9.1 8.7 7.7 8.1 9 9.5 7.5 ...
## $ Fuel.Consumption.Comb..L.100.km.: num 8.5 9.6 5.9 11.1 10.6 10 10.1 11.1 11.6 9.2 ...
## $ Fuel.Consumption.Comb..mpg. : int 33 29 48 25 27 28 28 25 24 31 ...
## $ CO2.Emissions.g.km. : int 196 221 136 255 244 230 232 255 267 212 ...
```

```
unique(d1$Vehicle.Class)
```

```
## [1] "COMPACT" "SUV - SMALL"
## [3] "MID-SIZE" "TWO-SEATER"
## [5] "MINICOMPACT" "SUBCOMPACT"
## [7] "FULL-SIZE" "STATION WAGON - SMALL"
## [9] "SUV - STANDARD" "VAN - CARGO"
## [11] "VAN - PASSENGER" "PICKUP TRUCK - STANDARD"
## [13] "MINIVAN" "SPECIAL PURPOSE VEHICLE"
## [15] "STATION WAGON - MID-SIZE" "PICKUP TRUCK - SMALL"
```


Data Preprocessing

High correlations between features, could cause collinearity issues

```
cor(d1$Engine.Size.L., d1$Cylinders)
```

```
## [1] 0.9276529
```

```
ggplot(d1, aes(x = Engine.Size.L., y = Cylinders)) +
  geom_point() +
  ggtitle("Scatterplot of Engine Size (L) and Number of Cylinders") +
  xlab("Engine Size (L)") +
  ylab("Cylinders")
```




```
cor(d1$Engine.Size.L.,d1$CO2.Emissions.g.km.)
```

```
## [1] 0.8511446
```

```
ggplot(d1, aes(x = Engine.Size.L., y = CO2.Emissions.g.km.)) +
  geom_point() +
  ggtitle("Scatterplot of Engine Size and CO2 Emissions ") +
  xlab("Engine Size (L)") +
  ylab("CO2 Emissions (g/km)")
```

Scatterplot of Engine Size and CO2 Emissions




```
cor(d1$Cylinders, d1$CO2.Emissions.g.km.)
```

```
## [1] 0.8326436
```

```
ggplot(d1, aes(x = Cylinders, y = CO2.Emissions.g.km.)) +  
  geom_point() +  
  ggtitle("Scatterplot of Cylinder and CO2 Emissions") +  
  xlab("Number of Cylinders") +  
  ylab("CO2 Emissions (g/km)")
```

Scatterplot of Cylinder and CO2 Emissions



Since I want to avoid multicollinearity and engine size has a higher correlation with the target variable, I chose to drop the cylinders feature.

```
d1 <- d1 %>% select(-Cylinders)
```

```
unique(d1$Transmission)
```

```
## [1] "AS5" "M6" "AV7" "AS6" "AM6" "A6" "AM7" "AV8" "AS8" "A7"
## [11] "A8" "M7" "A4" "M5" "AV" "A5" "AS7" "A9" "AS9" "AV6"
## [21] "AS4" "AM5" "AM8" "AM9" "AS10" "A10" "AV10"
```

Since there were too many models and it would be too difficult to encode them, I decided to drop the feature.

```
d1 <- d1 %>% select(-Model)
```

```
colnames(d1)
```

```
## [1] "Make" "Vehicle.Class"
## [3] "Engine.Size.L." "Transmission"
## [5] "Fuel.Type" "Fuel.Consumption.City..L.100.km."
## [7] "Fuel.Consumption.Hwy..L.100.km." "Fuel.Consumption.Comb..L.100.km."
## [9] "Fuel.Consumption.Comb..mpg." "CO2.Emissions.g.km."
```

```
unique(d1$Fuel.Type)
```

```
## [1] "Z" "D" "X" "E" "N"
```

Drop all types of gas mileage except for miles per gallon, as this is the industry norm.

```
d1 <- d1 %>% select(-Fuel.Consumption.City..L.100.km., -Fuel.Consumption.Hwy..L.100.km., -Fuel.Consumpt.
```

```
colnames(d1)
```

```
## [1] "Make" "Vehicle.Class"
## [3] "Engine.Size.L." "Transmission"
## [5] "Fuel.Type" "Fuel.Consumption.Comb..mpg."
## [7] "CO2.Emissions.g.km."
```

```
unique(d1$Make)
```

```
## [1] "ACURA" "ALFA ROMEO" "ASTON MARTIN" "AUDI"
## [5] "BENTLEY" "BMW" "BUICK" "CADILLAC"
## [9] "CHEVROLET" "CHRYSLER" "DODGE" "FIAT"
## [13] "FORD" "GMC" "HONDA" "HYUNDAI"
## [17] "INFINITI" "JAGUAR" "JEEP" "KIA"
## [21] "LAMBORGHINI" "LAND ROVER" "LEXUS" "LINCOLN"
## [25] "MASERATI" "MAZDA" "MERCEDES-BENZ" "MINI"
## [29] "MITSUBISHI" "NISSAN" "PORSCHE" "RAM"
## [33] "ROLLS-ROYCE" "SCION" "SMART" "SRT"
## [37] "SUBARU" "TOYOTA" "VOLKSWAGEN" "VOLVO"
## [41] "GENESIS" "BUGATTI"
```

Separate all car makes into either economy or luxury brands. A value of 0 is given to all economy cars and a 1 to all luxury brands. Add this column to the d1 dataframe.

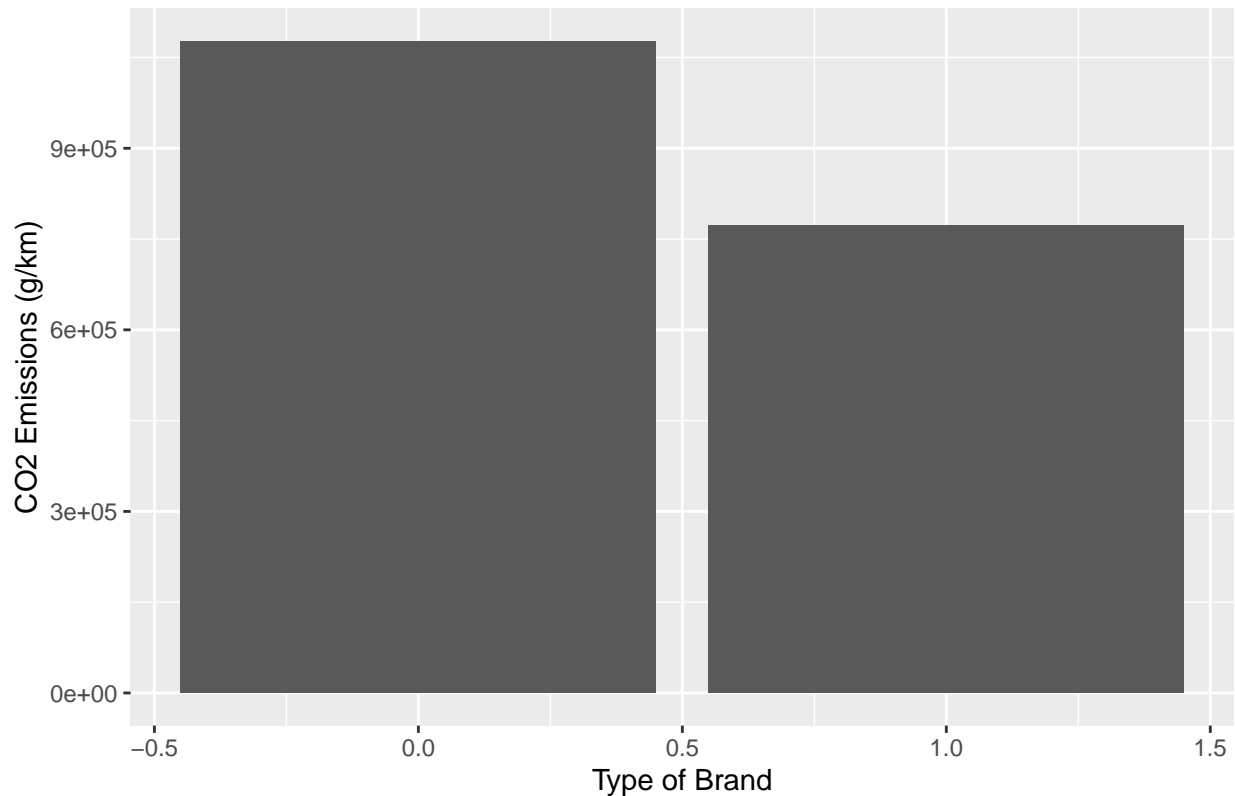
```
economy_brands <- c("BUICK", "CHEVROLET", "CHRYSLER", "DODGE", "FIAT", "FORD", "GMC", "HONDA", "HYUNDAI",
```

```
luxury_brands <- c("ACURA", "ALFA ROMEO", "ASTON MARTIN", "AUDI", "BENTLEY", "BMW", "CADILLAC", "CADILLAC",
```

```
d1$brand_encoded <- ifelse(d1$Make %in% economy_brands, 0, 1)
```

```
ggplot(d1, aes(x = brand_encoded, y = CO2.Emissions.g.km.)) +
  geom_col() +
  ggtitle("CO2 by Type of Make") +
  xlab("Type of Brand") +
  ylab("CO2 Emissions (g/km)")
```

CO2 by Type of Make



```
unique(d1$Transmission)
```

```
## [1] "AS5" "M6" "AV7" "AS6" "AM6" "A6" "AM7" "AV8" "AS8" "A7"
## [11] "A8" "M7" "A4" "M5" "AV" "A5" "AS7" "A9" "AS9" "AV6"
## [21] "AS4" "AM5" "AM8" "AM9" "AS10" "A10" "AV10"
```

To encode the different types of transmissions, I separated them into the following columns: Automatic sequential, automatic, manual, automated manual, and continuously variable transmission. I then encoded these and added them to the d1 dataframe.


```
Automatic_seq <- c("AS5", "AS6", "AS8", "AS7", "AS9", "AS4", "AS10")
Automatic <- c("A6", "A7", "A8", "A4", "A5", "A9")
Manual <- c("M6", "M7", "M5")
Automated_Manual <- c("AM6", "AM7", "AM5", "AM8", "AM9")

d1$Transmission_encoded <- ifelse(d1$Transmission %in% Automatic_seq, 1,
                                   ifelse(d1$Transmission %in% Automatic, 2,
                                           ifelse(d1$Transmission %in% Manual, 3,
                                                  ifelse(d1$Transmission %in% Automated_Manual, 4, 0))))
```

```
ggplot(d1, aes(x = Transmission_encoded, y = CO2.Emissions.g.km.)) +
  geom_col() +
  ggtitle("CO2 Emissions by Vehicle Transmission") +
  xlab("Type of Transmission")+
  ylab("CO2 Emissions (g/km)")
```



## CO2 Emissions by Vehicle Transmission



```
unique(d1$Vehicle.Class)
```

```
## [1] "COMPACT"           "SUV - SMALL"
## [3] "MID-SIZE"          "TWO-SEATER"
## [5] "MINICOMPACT"       "SUBCOMPACT"
## [7] "FULL-SIZE"         "STATION WAGON - SMALL"
## [9] "SUV - STANDARD"    "VAN - CARGO"
## [11] "VAN - PASSENGER"   "PICKUP TRUCK - STANDARD"
## [13] "MINIVAN"           "SPECIAL PURPOSE VEHICLE"
## [15] "STATION WAGON - MID-SIZE" "PICKUP TRUCK - SMALL"
```

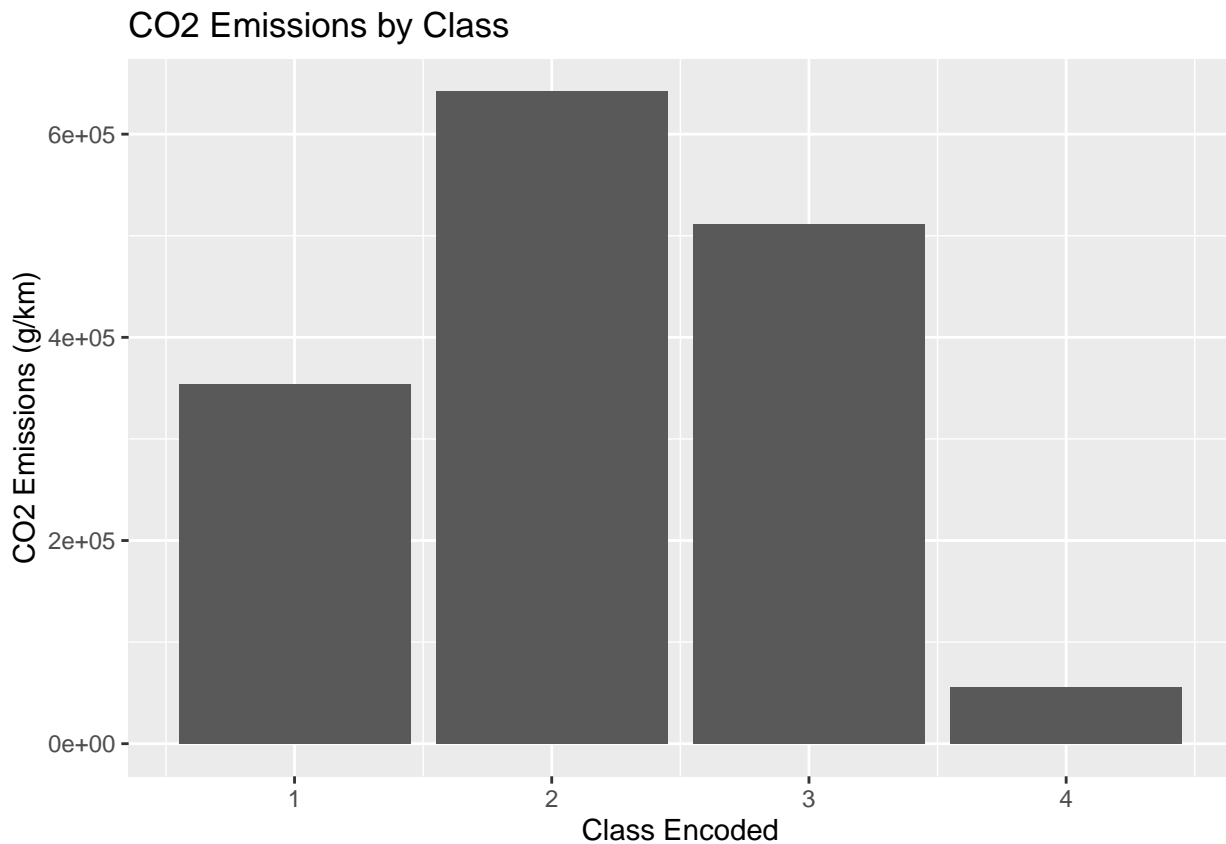
I separated the cars into coupes, sedans, SUVs, vans and trucks as these are the most common types of cars. I then encoded these and added them to the d1 dataframe.

```
coupe <- c("TWO-SEATER", "MINICOMPACT", "SUBCOMPACT")
sedan <- c("COMPACT", "MID-SIZE", "FULL-SIZE")
SUV <- c("SUV - SMALL", "SUV - STANDARD")
VAN <- c("VAN - CARGO", "VAN - PASSENGER", "MINIVAN")
Truck <- c("PICKUP TRUCK - SMALL", "PICKUP TRUCK - STANDARD")

d1$Class_encoded <- ifelse(d1$Vehicle.Class %in% coupe, 1,
                           ifelse(d1$Vehicle.Class %in% sedan, 2,
                                   ifelse(d1$Vehicle.Class %in% SUV, 3,
                                         ifelse(d1$Vehicle.Class %in% VAN, 4,
                                                ifelse(d1$Transmission %in% Truck, 5, 0))))))
```

```
d1 <- subset(d1, Class_encoded != 0)
```

```
ggplot(d1, aes(x = Class_encoded, y = CO2.Emissions.g.km.)) +  
  geom_col() +  
  ggtitle("CO2 Emissions by Class") +  
  xlab("Class Encoded") +  
  ylab("CO2 Emissions (g/km)")
```




I changed the fuel type from a string to a factor variable to encode it. The fuel had many types including normal gas, premium gas, no gas (electric), diesel, and ethanol.

```
d1$Fuel.Type <- as.factor(d1$Fuel.Type)
```

```
ggplot(d1, aes(x=Fuel.Type, y = CO2.Emissions.g.km.)) +  
  geom_col() +  
  ggtitle("CO2 Emissions by Fuel Type") +  
  xlab("Fuel Type") +  
  ylab("CO2 Emissions (g/km)")
```

## CO2 Emissions by Fuel Type



Fitting linear model

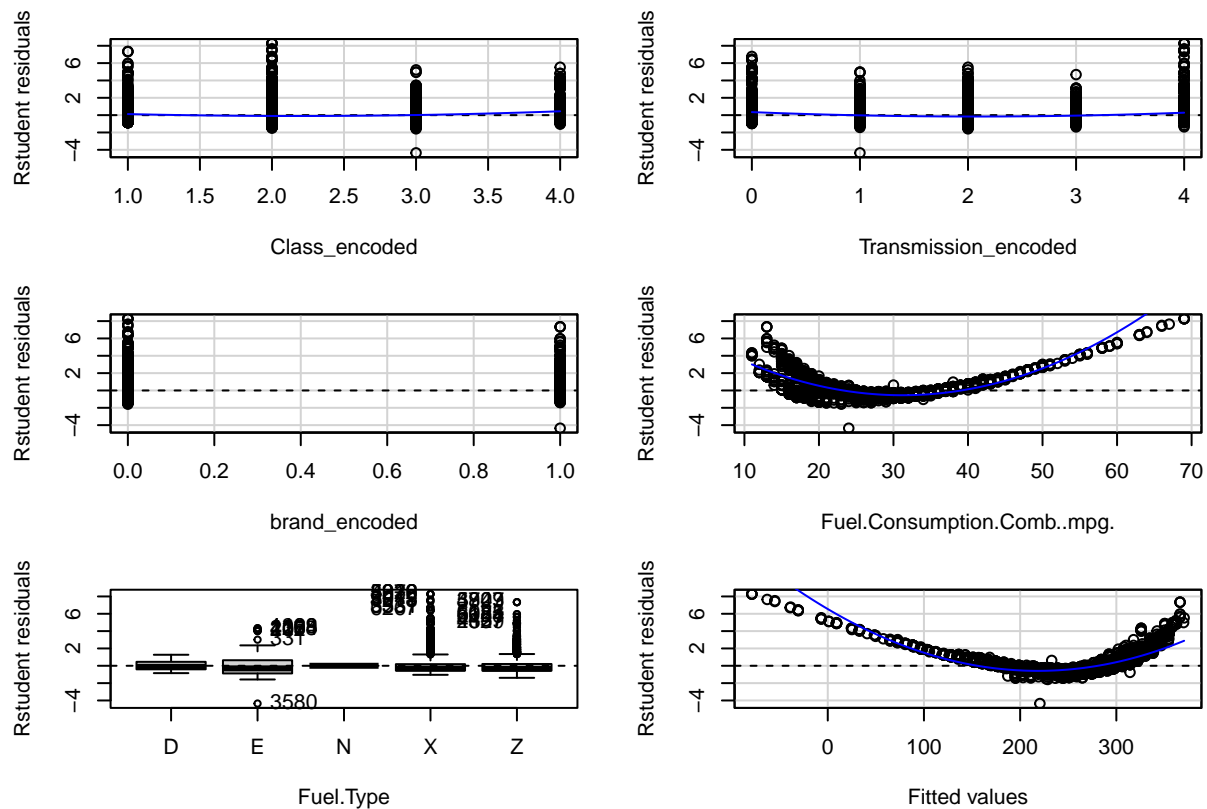
```
m1 <- lm(CO2.Emissions.g.km. ~ Class_encoded + Transmission_encoded + brand_encoded + Fuel.Consumption.Comb..mpg. + Fuel.Type, data = d1)
```

```
summary(m1)
```

```
##
## Call:
## lm(formula = CO2.Emissions.g.km. ~ Class_encoded + Transmission_encoded +
##     brand_encoded + Fuel.Consumption.Comb..mpg. + Fuel.Type,
##     data = d1)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -92.533 -12.090  -6.842   4.767 174.884
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    498.17024     3.23313   154.083 < 2e-16 ***
## Class_encoded      1.67897     0.42519     3.949 7.94e-05 ***
## Transmission_encoded -0.11372     0.25725    -0.442  0.658
## brand_encoded    -0.26713     0.76948    -0.347  0.728
## Fuel.Consumption.Comb..mpg. -7.96035     0.04574 -174.052 < 2e-16 ***
## Fuel.TypeE      -91.24529     2.63596   -34.616 < 2e-16 ***
## Fuel.TypeN     -113.28668    21.46294    -5.278 1.35e-07 ***
## Fuel.TypeX      -30.69284     2.15742   -14.227 < 2e-16 ***
```

```
## Fuel.TypeZ          -29.42117    2.17885   -13.503   < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 21.34 on 6297 degrees of freedom
## Multiple R-squared:  0.8695, Adjusted R-squared:  0.8694
## F-statistic: 5245 on 8 and 6297 DF,  p-value: < 2.2e-16
```

```
residualPlots(m1, type = "rstudent")
```




```
##              Test stat Pr(>|Test stat|)
## Class_encoded      9.4850      <2e-16 ***
## Transmission_encoded 13.6601      <2e-16 ***
## brand_encoded       0.9395      0.3475
## Fuel.Consumption.Comb..mpg. 161.7271      <2e-16 ***
## Fuel.Type
## Tukey test         152.8104      <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Since the graph displays heavy curvature, a linear model does not seem to be appropriate here.

```
studentized_residuals <- rstudent(m1)
hist(studentized_residuals)
```

## Histogram of studentized\_residuals



```
model_rf <- randomForest(d1$CO2.Emissions.g.km. ~ Class_encoded + Transmission_encoded + brand_encoded)
summary(model_rf)
```

```
##               Length Class  Mode
## call           4      -none- call
## type           1      -none- character
## predicted      6306     -none- numeric
## mse            500     -none- numeric
## rsq            500     -none- numeric
## oob.times      6306     -none- numeric
## importance       5     -none- numeric
## importanceSD     0     -none-  NULL
## localImportance  0     -none-  NULL
## proximity        0     -none-  NULL
## ntree           1     -none- numeric
## mtry            1     -none- numeric
## forest          11     -none- list
## coefs           0     -none-  NULL
## y              6306     -none- numeric
## test            0     -none-  NULL
## inbag           0     -none-  NULL
## terms           3      terms  call
```

```
predictions <- predict(model_rf, d1)
y_actual <- d1$CO2.Emissions.g.km.
rss <- sum((y_actual - predictions)^2)
tss <- sum((y_actual - mean(y_actual))^2)
print(1 - (rss / tss))
```

```
## [1] 0.8068461
```

The  $r^2$  value that we are getting here is 81%. This is a decent value, but we can try to enhance it with hyperparameter tuning.

```
install.packages("randomForest")
```

```
##  
## The downloaded binary packages are in  
## /var/folders/c3/2w85_lw50qdgkzcnmbnz057m0000gn/T//Rtmp4eGXuX/downloaded_packages
```

```
install.packages("caret")
```

```
##  
## The downloaded binary packages are in  
## /var/folders/c3/2w85_lw50qdgkzcnmbnz057m0000gn/T//Rtmp4eGXuX/downloaded_packages
```

```
library(randomForest)  
library(caret)
```

```
## Loading required package: lattice
```

```
set.seed(123)  
  
# Define a tuning grid for mtry  
tune_grid <- expand.grid(  
  mtry = seq(2, 6, by = 1)  
)
```

Using grid search to find best possible parameters for random forest model.

```
control <- trainControl(  
  method = "cv",  
  number = 5,  
  search = "grid"  
)
```

Using the caret package, we can use the optimal parameters to

```
tuned_model <- train(  
  CO2.Emissions.g.km. ~ Class_encoded + Transmission_encoded + brand_encoded + Fuel.Consumption.Comb...mpg,   
  data = d1,  
  method = "rf",  
  tuneGrid = tune_grid,  
  trControl = control,  
  ntree = 500  
)  
  
print(tuned_model$bestTune)
```

```
##      mtry  
## 5      6
```

```
print(tuned_model)
```

```
## Random Forest
##
## 6306 samples
##    5 predictor
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 5045, 5046, 5043, 5045, 5045
## Resampling results across tuning parameters:
##
##  mtry  RMSE      Rsquared  MAE
##  2     21.812118  0.9293333  14.824141
##  3     10.847032  0.9786125   7.049744
##  4      6.007994  0.9909151   4.027598
##  5      4.768108  0.9935391   3.327970
##  6      4.558736  0.9940332   3.204856
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was mtry = 6.
```